

RDF/OWL and SPARQL instead of NoSQL databases

Philip Cannata (Oracle and the University of Texas), Isabella Bhardwaj (University of Texas), Michael Porras (Rackspace), Steven Spohrer (University of Texas and Samsung Research), Nigel Jacobs (Oracle), Andrew Oldag, Francisco Garcia

1. Introduction

NoSQL database systems, especially MongoDB¹ (MongoDB, 2014), have been gaining notoriety over the past several years purportedly for several reason 1) their data model (e.g., JSON) and its “schema-less flexibility” is a better match for data-driven web applications; 2) they are well suited for Database-as-a-Service (DBaaS) implementations (mongolab, 2014); and 3) they can scale across small, inexpensive systems, i.e., they can “scale-out”. However, the underlying “schema-less flexible” data model for these systems is very similar to technology that has a very long history in the form of Entity-Attribute-Value (EAV) database systems (Stead, Hammond, & Straube, 1982) (McDonald, Blevins, Tierney, & Martin, 1988).

EAV technology has evolved into the standards-based RDF/OWL and SPARQL technology of today and this paper will show that RDF/OWL and SPARQL are practical alternatives to current NoSQL databases for building data-driven web applications and for DBaaS implementations. (Scale-out won’t be discussed in this paper because it is believed to be orthogonal to the discussion.)

As a case study to demonstrate the capabilities of RDF/OWL and SPARQL, this paper will present an implementation of a simple, flask-based (Flask, 2014) “book” website developed using the MongoDB NoSQL database system and the same website developed with RDF/OWL and SPARQL embedded in a system called ReL², This paper will also show how to build DBaaS applications using RESTful ReL. Both

¹ Similarities with the Cassandra (Cassandra, 2014) NoSQL system will also be shown in Section 3 of this paper.

² ReL (Relation Language) is a python-based, data management system that uses RDF/OWL and SPARQL as its tuple manager. In addition to supporting RDF/OWL and SPARQL, ReL is data model agnostic and allows data manipulation and retrieval using a mix and match of many different higher-level data models, including the Relational Model, a Semantic Model based upon the work of Hammer and McLeod (Hammer & McLeod, 1981), and the OO python model. The ReL Relational Model, which is automatically translated to RDF/OWL and SPARQL, is used in this paper. However, because ReL is a python-based system, it’s trivial to also support JSON by translating JSON into one of the other supported data models and to translate results back into JSON. This has been done in several ReL applications.

of these examples will show that RDF/OWL and SPARQL are very well suited for these types of applications.

These examples also demonstrate that a major weakness of NoSQL databases – that there is no universal query language - can be overcome by standardizing on the SPARQL query language.

This paper also shows that, given a suitable implementation of an RDF triple-store (e.g. (Oracle Graph, 2014)), support for standard transaction capabilities (i.e., read committed, serializable, and ACID) can be made available to NoSQL-like applications. This addresses another major weakness of NoSQL databases; that they do not provide this level of transaction support. MongoDB, for example, only supports single statement transactions and the entire Database³ is locked by the transaction statement.

Figure 1 shows the main menu for the “book” website example used in this paper.




Figure 1

2. Inserting data into the book website database

Figure 2 shows the menu for adding a book in this web application.




Figure 2

When the submit button is clicked, the following code is executed.

A. For the MongoDB application

³ A Mongo Database is a container for a set of Collections and a Collection is a container for a set of JSON Documents.

First, a MongoDB connection is established:

```
from flask import Flask, render_template, redirect, request
import pymongo

app = Flask(__name__, static_url_path = "")
connection_string = "mongodb://127.0.0.1"
connection = pymongo.MongoClient(connection_string)
database = connection.books
books = database.mybooks
```

Then, the insert is done using the MongoDB insert API:

```
@app.route('/add/', methods=['GET', 'POST'])
def add():
    if request.method == 'POST':
        new_data = {k : v for k, v in request.form.items()}
        books.insert(new_data)
        return render_template('add.html', alert =
            "success")
    else:
        return render_template('add.html', alert="")
```

B. For the ReL application

First, a connection to an Oracle RDF datastore is established:

```
from flask import Flask, render_template, redirect, request

app = Flask(__name__, static_url_path = "")
conn = connectTo 'jdbc:oracle:thin:@host:1521:orcl' 'user'
        'password' 'rdf_mode' 'bookApp'
```

Then, the insert is done using a standard SQL insert; however, no table named “books” was created beforehand, so this is “scheme-less,” and “flexible” just like MongoDB:

```
@app.route('/add/', methods=['GET', 'POST'])
def add():
    if request.method == 'POST':
        new_data = {k : v for k, v in request.form.items()}
        values = (str(new_data['title']),
            str(new_data['author']), str(new_data['genre']),
            str(new_data['description']))
        SQL on conn """insert into books(title, author,
            genre, description) values"""values
        return render_template('add.html', alert =
```

```

        "success")
    else:
        return render_template('add.html', alert="")

```

Behind the scenes, ReL converts the SQL insert into a series of several RDF/OWL insert statements as follows⁴:

```

BEGIN
commit ;
set transaction isolation level serializable5 ;
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#89',
'owl#title', '"A Profile of Mathematical
Logic"^^xsd:string'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdf:type', 'owl:DatatypeProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdfs:domain', 'owl#books'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdf:range', 'rdfs:xsd:string'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdf:type', 'owl:FunctionalProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#89',
'owl#author', '"Howard DeLong"^^xsd:string'));

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#author',
'rdf:type', 'owl:DatatypeProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#author',
'rdfs:domain', 'owl#books'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#author',
'rdf:range', 'rdfs:xsd:string'));

```

⁴ URIs have been abbreviated to help with readability.

⁵ Notice that Oracle allows RDF triple-store statements to be wrapped in standard SQL. This means standard transaction processing can be done with RDF triple-store statements. The same is true for Oracle SPARQL statements as will be seen in Section 4.

```

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#author',
'rdf:type', 'owl:FunctionalProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#89',
'owl#genre', '"Math"^^xsd:string'));

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#genre',
'rdf:type', 'owl:DatatypeProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#genre',
'rdfs:domain', 'owl#books'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#genre',
'rdf:range', 'rdfs:xsd:string'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#genre',
'rdf:type', 'owl:FunctionalProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#89',
'owl#description', '"The best"^^xsd:string'));

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>',
'owl#description', 'rdf:type', 'owl:DatatypeProperty'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>',
'owl#description', 'rdfs:domain', 'owl#books'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>',
'owl#description', 'rdf:range', 'rdfs:xsd:string'));
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>',
'owl#description', 'rdf:type', 'owl:FunctionalProperty'));
END ;
/

```

Notice that some of the INSERT statements above insert appropriate OWL schema information into the RDF tuple-store, e.g., for the “title” attribute, the following is inserted:

```

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdf:type', 'owl:DatatypeProperty'))

```

```

INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SEQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdfs:domain', 'owl#books'))
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SEQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdfs:range', 'rdfs:xsd:string'))
INSERT INTO BOOK_DATA VALUES ( BOOK_APP_SEQNC.nextval,
SDO_RDF_TRIPLE_S('FALL2014_CS347_PROF:<owl>', 'owl#title',
'rdfs:type', 'owl:FunctionalProperty'))

```

This schema information is used when querying the book database.

3. Searching the book website database

Figure 3 shows the menu for searching for a book in this web application.



The image shows a web application interface for searching a database. It has a large heading "Search The Database:" in a bold, sans-serif font. Below the heading is a search form with a text input field containing the text "Howard DeLong" and a button labeled "Submit". The input field has a light yellow background, and the button is a simple rectangular button.

Figure 3

When the submit button is clicked, the following code is executed:

A. For the MongoDB application

The search is done using the MongoDB “find” API:

```

@app.route('/search/', methods=['GET', 'POST'])
def search():
    if request.method == 'POST':
        query = request.form['query']
        return render_template('search.html', posting=True,
                               query=query, title_results = books.find({
                                   'title':query}),author_results = books.find({
                                   'author':query}))
    else:
        return render_template('search.html', posting =
                               False)

```

B. For the ReL application

The ReL search is done using standard SQL select statements as follows (the SQL select statements return a python tuple of tuples):

```

@app.route('/search/', methods=['GET', 'POST'])
def search():
    if request.method == 'POST':
        query = request.form['query']

        titles6 = SQL on conn """select title, author from
            books where title = '""query""'"""
        authors = SQL on conn """select title, author from
            books where author = '""query""'"""

        title_dict = convert_to_dict(titles)
        author_dict = convert_to_dict(authors)
        genre_dict = {}

        no_results = title_dict == 0 and author_dict == 0
            and genre_dict == 0
        return render_template('search.html', posting=True,
            query=query, no_results=no_results,
            title_results=title_dict,
            author_results=author_dict,
            genre_results=genre_dict)

    else:
        return render_template('search.html',
            posting=False)

```

Behind the scenes, ReL converts the SQL select statements into SPARQL statements, as shown below^{7,8}:

```

SELECT v1 "title", v2 "author"
FROM TABLE(SEM_MATCH('SELECT * WHERE {
    ?s1 rdf:type :books .
    OPTIONAL { ?s1 :title ?v1 }
    OPTIONAL { ?s1 :author ?v2 }
    ?s1 :title ?f1 .
    FILTER(?f1 = "Howard DeLong") }' ,

```

⁶ The SQL statement returns a python tuple of tuples.

⁷ Notice, the SPARQL statement is in an Oracle Table function that is a part of a standard SQL statement.

⁸ The use of the OPTIONAL pattern in the SPARQL statements means that each of the attributes modified by the OPTIONAL pattern will optionally be part of a returned tuple (row). This is like the CASSANDRA NoSQL database system where “Unlike a table in an RDBMS, different rows in the same column family do not have to share the same set of columns, and a column may be added to one or multiple rows at any time.” (Cassandra, 2014)

```
SEM_MODELS('FALL2014_CS347_PROF'), null,
SEM_ALIASES( SEM_ALIAS('',
'http://www.example.org/people.owl#')), null) )
```

```
SELECT v1 "title", v2 "author"
FROM TABLE(SEM_MATCH('SELECT * WHERE {
    ?s1 rdf:type :books .
    OPTIONAL { ?s1 :title ?v1 }
    OPTIONAL { ?s1 :author ?v2 }
    ?s1 :author ?f1 .
    FILTER(?f1 = "Howard DeLong") }' ,
SEM_MODELS('FALL2014_CS347_PROF'), null,
SEM_ALIASES( SEM_ALIAS('',
'http://www.example.org/people.owl#')), null) )
```

4. Transaction support in ReL

As discussed in a previous footnote, the SPARQL statements shown above can be included in a regular Oracle transaction (see the example below). The same is true for all other database operations in ReL. So, ReL can provide traditional transaction support (i.e., read committed, serializable, and ACID) for all of its database operations. (Notice in Section 2 above, the INSERTS were wrapped in PL/SQL BEGIN and END statements, which made the INSERTS ACID). One of the major criticisms of NoSQL databases is that they don't provide traditional transaction support. ReL does not suffer from this problem when using Oracle's implementation of RDF/OWL and SPARQL.

```
commit ;
```

```
set transaction isolation level serializable ;
```

```
SELECT v1 "title", v2 "author"
FROM TABLE(SEM_MATCH('SELECT * WHERE {
    ?s1 rdf:type :books .
    OPTIONAL { ?s1 :title ?v1 }
    OPTIONAL { ?s1 :author ?v2 }
    ?s1 :title ?f1 .
    FILTER(?f1 = "Howard DeLong") }' ,
SEM_MODELS('FALL2014_CS347_PROF'), null,
SEM_ALIASES( SEM_ALIAS('',
'http://www.example.org/people.owl#')), null) );
```

```
SELECT v1 "title", v2 "author"
FROM TABLE(SEM_MATCH('SELECT * WHERE {
    ?s1 rdf:type :books .
    OPTIONAL { ?s1 :title ?v1 }
    OPTIONAL { ?s1 :author ?v2 }
    ?s1 :author ?f1 .
```



```

    FILTER(?f1 = "Howard DeLong") }' ,
SEM_MODELS('FALL2014_CS347_PROF'), null,
SEM_ALIASES( SEM_ALIAS('',
'http://www.example.org/people.owl#')), null) );

```

5. DBaaS - RESTful ReL

ReL can also be run as a RESTful (Fielding, 2000) server, which means ReL SQL calls (e.g., ""select title, author from books where title = ""query""""""""", which was discussed above) can be embedded in any environment that supports the CURL function. For instance, we use RESTful ReL in R (Project, 2014) to access data and convert it to R data frames for analysis. The same “data model-to-RDF/OWL and SPARQL” translations that were discussed above can be used with Restful ReL. Here’s how ReL can be invoked from R to query the standard Oracle emp table:

```

d = getURL( URLEncode('host:5000/rest/native/?query =
"select * from emp"), httpheader =
c(DB='jdbc:oracle:thin:@host:1521:orcl', USER='user',
PASS='password', MODE='rdf_mode', MODEL='Fall2014',
returnFor = 'R'), verbose = TRUE)

```

The “returnFor = 'R'” httpheader parameter value above directs RESTful ReL to return data in the following format:

```

d
"list(c('COMM', 'HIREDATE', 'JOB', 'DEPTNO', 'SAL',
'ENAME', 'MGR', 'EMPNO'), list(c('NULL', 1400, 'NULL',
'NULL', 500, 'NULL', 300, 'NULL', 'NULL', 'NULL', 'NULL',
'NULL', 'NULL', 'NULL'),c('23-JAN-1982', '28-SEP-1981', '1-
MAY-1981', '3-DEC-1981', '22-FEB-1981', '9-JUN-1981', '20-
FEB-1981', '8-SEP-1981', '12-JAN-1983', '09-DEC-1982', '17-
NOV-1981', '17-DEC-1980', '3-DEC-1981', '2-APR-
1981'),c('CLERK', 'SALESMAN', 'MANAGER', 'ANALYST',
'SALESMAN', 'MANAGER', 'SALESMAN', 'SALESMAN', 'CLERK',
'ANALYST', 'PRESIDENT', 'CLERK', 'CLERK', 'MANAGER'),c(10,
30, 30, 20, 30, 10, 30, 30, 20, 20, 10, 20, 30, 20),c(1300,
1250, 2850, 3000, 1250, 2450, 1600, 1500, 1100, 3000, 5000,
800, 950, 2975),c('MILLER', 'MARTIN', 'BLAKE', 'FORD',
'WARD', 'CLARK', 'ALLEN', 'TURNER', 'ADAMS', 'SCOTT',
'KING', 'SMITH', 'JAMES', 'JONES'),c(7782, 7698, 7839,
7566, 7698, 7839, 7698, 7698, 7788, 7566, 'NULL', 7902,
7698, 7839),c(7934, 7654, 7698, 7902, 7521, 7782, 7499,
7844, 7876, 7788, 7839, 7369, 7900, 7566)))"

```

Then, the data can be converted to an R data frame using the following two commands:

```
df <- data.frame(eval(parse(text=substring(d,1))) [2])

colnames(df) <- unlist(eval(parse(text=substring(d,1))) [1])
```

Finally, `head(df)` results in the following:

```
head(df)
```

	COMM	HIREDATE	JOB	DEPTNO	SAL	ENAME	MGR	EMPNO
1	NULL	23-JAN-1982	CLERK	10	1300	MILLER	7782	7934
2	1400	28-SEP-1981	SALESMAN	30	1250	MARTIN	7698	7654
3	NULL	1-MAY-1981	MANAGER	30	2850	BLAKE	7839	7698
4	NULL	3-DEC-1981	ANALYST	20	3000	FORD	7566	7902
5	500	22-FEB-1981	SALESMAN	30	1250	WARD	7698	7521
6	NULL	9-JUN-1981	MANAGER	10	2450	CLARK	7839	7782

6. Summary

In an InfoWorld article (Oliver, 2014), the author claims, “the time for NoSQL standards is now”. But, RDF/OWL and SPARQL are standards that exist now and this paper has demonstrated that they are perfectly well suited for building “schema-less”, flexible NoSQL-type applications and DBaaS applications. What’s more, unlike other NoSQL systems, RDF/OWL and SPARQL systems like ReL can support standard, read committed, serializable, and ACID transaction processing. So, maybe “the time for NoSQL to use existing standards is now”.

The Oracle implementation of RDF/OWL and SPARQL (Oracle Graph, 2014) was used for the applications in this paper; however, any proper implementation of these standards could be used instead.

This paper did not discuss the “scale-out rather than scale-up” proposition of NoSQL databases, but that debate has a decades long history and needs no more discussion here. However, there is no reason that RDF/OWL and SPARQL could not be used just as effectively in a “scale-out” system. As a matter of fact, ReL and RESTful ReL would run blazingly fast in Oracle’s Exadata (Exadata, 2014) environment, which is effectively “scale-out” albeit not inexpensive.

7. Bibliography

Cassandra. (2014). Retrieved from http://en.wikipedia.org/wiki/Apache_Cassandra

Exadata. (2014). Retrieved from <https://www.oracle.com/engineered-systems/exadata/index.html>

Fielding, R. (2000). *REST*. Retrieved from http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm

Flask. (2014). Retrieved from Flask web development, one drop at a time: <http://flask.pocoo.org/>

Hammer, M., & McLeod, D. (1981). Database Description with SDM: A Semantic Database Model. *ACM Transaction of Database* , 6 (3).

McDonald, C., Blevins, L., Tierney, W., & Martin, D. (1988). The Regenstrief Medical Records. *MD Computing* (5(5)) , 34-47.

MongoDB. (2014). Retrieved from MongoDB: <http://www.mongodb.org>

mongolab. (2014). Retrieved from <https://mongolab.com/>

Oliver, A. (2014). *The time for NoSQL standards is now*. Retrieved from InfoWorld: <http://www.infoworld.com/article/2615807/nosql/the-time-for-nosql-standards-is-now.html>

Oracle Graph. (2014). Retrieved from <http://www.oracle.com/us/products/database/options/spatial/overview/index.html>

Project, R. (2014). Retrieved from <http://www.r-project.org/>

Stead, W., Hammond, W., & Straube, M. (1982, November). A Chartless Record—Is It Adequate? *Proceedings of the Annual Symposium on Computer Application in Medical Care* , 89-94.